

```
In [4]: ## Importing Libraries

import numpy as np # Linear algebra operations
import pandas as pd # Data processing and analysis
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
In [46]: ## Upload dataset

df = pd.read_csv('/Users/serenaygoler/heart disease.csv')

df.head() # Displays the first 5 rows.
```

```
Out[46]:
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR
0	40	M	ATA	140	289	0	Normal	172
1	49	F	NAP	160	180	0	Normal	156
2	37	M	ATA	130	283	0	ST	98
3	48	F	ASY	138	214	0	Normal	108
4	54	M	NAP	150	195	0	Normal	122

```
In [48]: df.tail() # Display the last 5 rows.
```

```
Out[48]:
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxH
913	45	M	TA	110	264	0	Normal	13
914	68	M	ASY	144	193	1	Normal	14
915	57	M	ASY	130	131	0	Normal	17
916	57	F	ATA	130	236	0	LVH	17
917	38	M	NAP	138	175	0	Normal	17

```
In [50]: df.info() # Prints name and type of variables, number of observations, and c
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                    918 non-null    int64
1   Sex                    918 non-null    object
2   ChestPainType          918 non-null    object
3   RestingBP              918 non-null    int64
4   Cholesterol            918 non-null    int64
5   FastingBS              918 non-null    int64
6   RestingECG            918 non-null    object
7   MaxHR                  918 non-null    int64
8   ExerciseAngina         918 non-null    object
9   Oldpeak                918 non-null    float64
10  ST_Slope               918 non-null    object
11  HeartDisease           918 non-null    int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
```

```
In [52]: df.shape # Displays the number of rows and columns in the dataset.
```

```
Out[52]: (918, 12)
```

```
In [54]: df.isna().sum() # Counts missing values in each column.
```

```
Out[54]: Age                0
Sex                  0
ChestPainType        0
RestingBP            0
Cholesterol          0
FastingBS            0
RestingECG           0
MaxHR                0
ExerciseAngina        0
Oldpeak              0
ST_Slope             0
HeartDisease         0
dtype: int64
```

```
In [56]: df.duplicated().sum() # Counts the number of duplicate rows.
```

```
Out[56]: 0
```

```
In [58]: ## Each subplot shows the frequency of categories for a specific column.
```

```
categorical_cols = ["HeartDisease", "Sex", "ChestPainType", "FastingBS", "RestingECG",
                    "ExerciseAngina", "ST_Slope"]

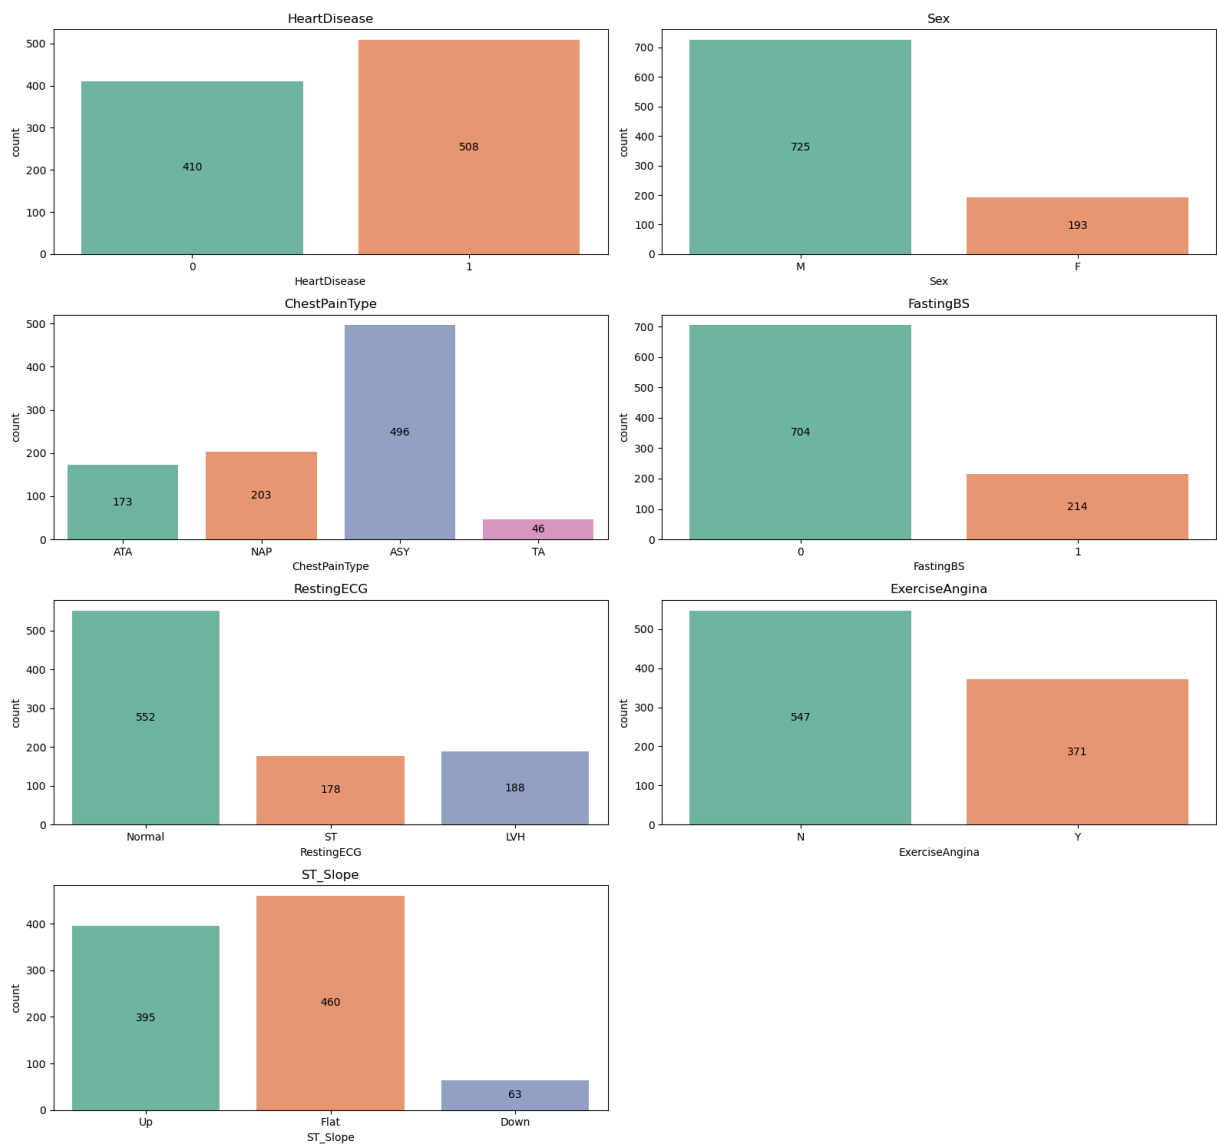
fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(16, 15))
axes = axes.flatten()

for ax, col in zip(axes, categorical_cols):
    sns.countplot(x=df[col], hue=df[col], ax=ax, palette="Set2", legend=False)
    ax.set_title(col)
    for container in ax.containers:
```

```
ax.bar_label(container, label_type="center")

for ax in axes[len(categorical_cols):]:
    fig.delaxes(ax)

plt.tight_layout()
plt.show()
```



```
In [60]: ## Provides summary statistics for numeric columns, rounded to 2 decimals and
df.describe().round(2).T
```

Out [60]:

	count	mean	std	min	25%	50%	75%	max
<b>Age</b>	918.0	53.51	9.43	28.0	47.00	54.0	60.0	77.0
<b>RestingBP</b>	918.0	132.40	18.51	0.0	120.00	130.0	140.0	200.0
<b>Cholesterol</b>	918.0	198.80	109.38	0.0	173.25	223.0	267.0	603.0
<b>FastingBS</b>	918.0	0.23	0.42	0.0	0.00	0.0	0.0	1.0
<b>MaxHR</b>	918.0	136.81	25.46	60.0	120.00	138.0	156.0	202.0
<b>Oldpeak</b>	918.0	0.89	1.07	-2.6	0.00	0.6	1.5	6.2
<b>HeartDisease</b>	918.0	0.55	0.50	0.0	0.00	1.0	1.0	1.0

```
In [62]: # Count how many Cholesterol values are zero
chol_zero_count = (df["Cholesterol"] == 0).sum()

# Count how many RestingBP values are zero
bp_zero_count = (df["RestingBP"] == 0).sum()

print(f"Number of Cholesterol values equal to 0: {chol_zero_count}")
print(f"Number of RestingBP values equal to 0: {bp_zero_count}")
```

Number of Cholesterol values equal to 0: 172

Number of RestingBP values equal to 0: 1

```
In [64]: ## Filters out rows where Cholesterol equals zero and returns summary statis

print(df[df["Cholesterol"] != 0]["Cholesterol"].describe())
```

```
count    746.000000
mean     244.635389
std       59.153524
min       85.000000
25%      207.250000
50%      237.000000
75%      275.000000
max      603.000000
Name: Cholesterol, dtype: float64
```

```
In [66]: # With zeros included
print("=== With Zero values Included ===")
print(df.groupby("HeartDisease")["Cholesterol"].describe())

# Zeros removed
print("\n=== With zero values removed ===")
print(df[df["Cholesterol"] != 0].groupby("HeartDisease")["Cholesterol"].desc
```

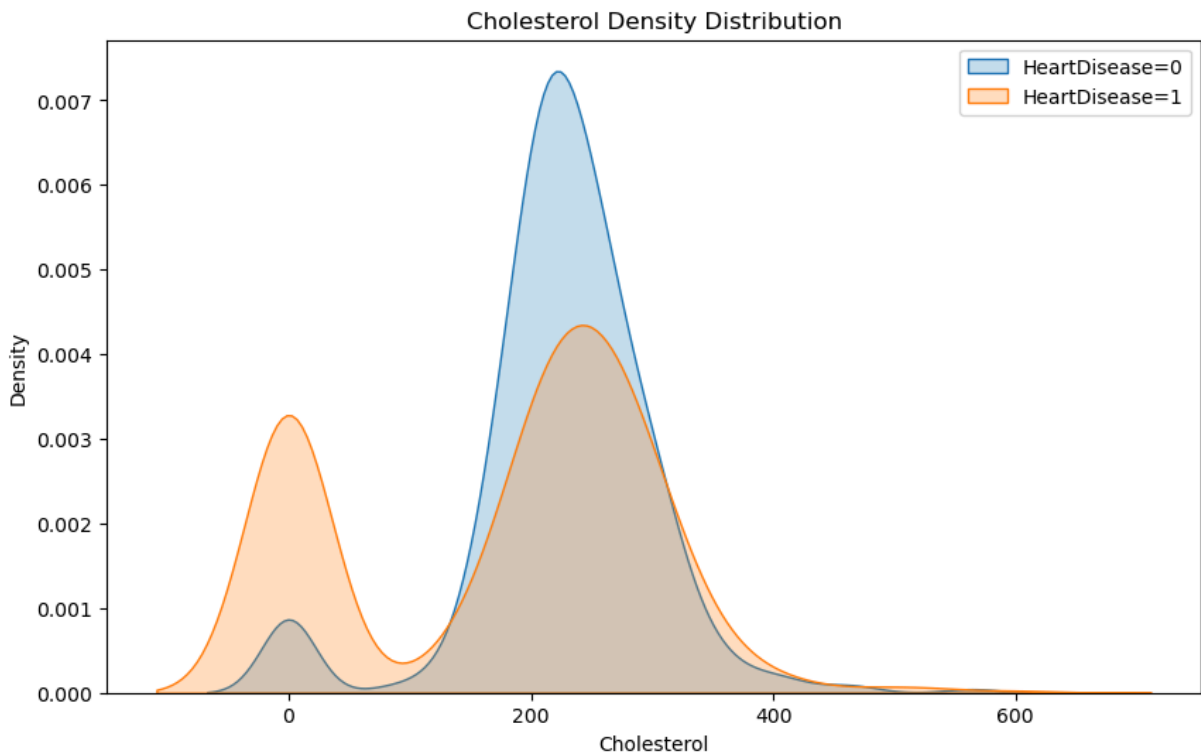
=== With Zero values Included ===

	count	mean	std	min	25%	50%	75%	m
ax								
HeartDisease								
0	410.0	227.121951	74.634659	0.0	197.25	227.0	266.75	56
4.0								
1	508.0	175.940945	126.391398	0.0	0.00	217.0	267.00	60
3.0								

=== With zero values removed ===

	count	mean	std	min	25%	50%	75%	m
ax								
HeartDisease								
0	390.0	238.769231	55.394617	85.0	203.0	231.5	269.00	56
4.0								
1	356.0	251.061798	62.462713	100.0	212.0	246.0	283.25	60
3.0								

```
In [68]: # Plot the cholesterol distribution for HeartDisease=0 and HeartDisease=1 us
plt.figure(figsize=(10,6))
sns.kdeplot(df[df["HeartDisease"]==0]["Cholesterol"], label="HeartDisease=0")
sns.kdeplot(df[df["HeartDisease"]==1]["Cholesterol"], label="HeartDisease=1")
plt.legend()
plt.title("Cholesterol Density Distribution")
plt.show()
```



```
In [70]: # This block cleans the dataset by:
# 1. Removing rows where RestingBP = 0 (unrealistic values).
# 2. Calculating group-wise medians of Cholesterol (by HeartDisease) excluding
# 3. Replacing Cholesterol values of zero with the corresponding group median
# 4. Checking that no zero values remain.
# 5. Displaying summary statistics of Cholesterol by HeartDisease after cleaning
```

```

df_clean = df.copy()
df_clean = df_clean[df_clean["RestingBP"] != 0].copy()

medians = (
    df_clean[df_clean["Cholesterol"] != 0]
    .groupby("HeartDisease")["Cholesterol"]
    .median()
)

mask_zero = df_clean["Cholesterol"] == 0
df_clean["Cholesterol"] = df_clean["Cholesterol"].astype(float)
df_clean.loc[mask_zero, "Cholesterol"] = (
    df_clean.loc[mask_zero, "HeartDisease"].map(medians)
)

print("Remaining zeros:", (df_clean["Cholesterol"] == 0).sum())
print(df_clean.groupby("HeartDisease")["Cholesterol"].describe())

```

```

Remaining zeros: 0

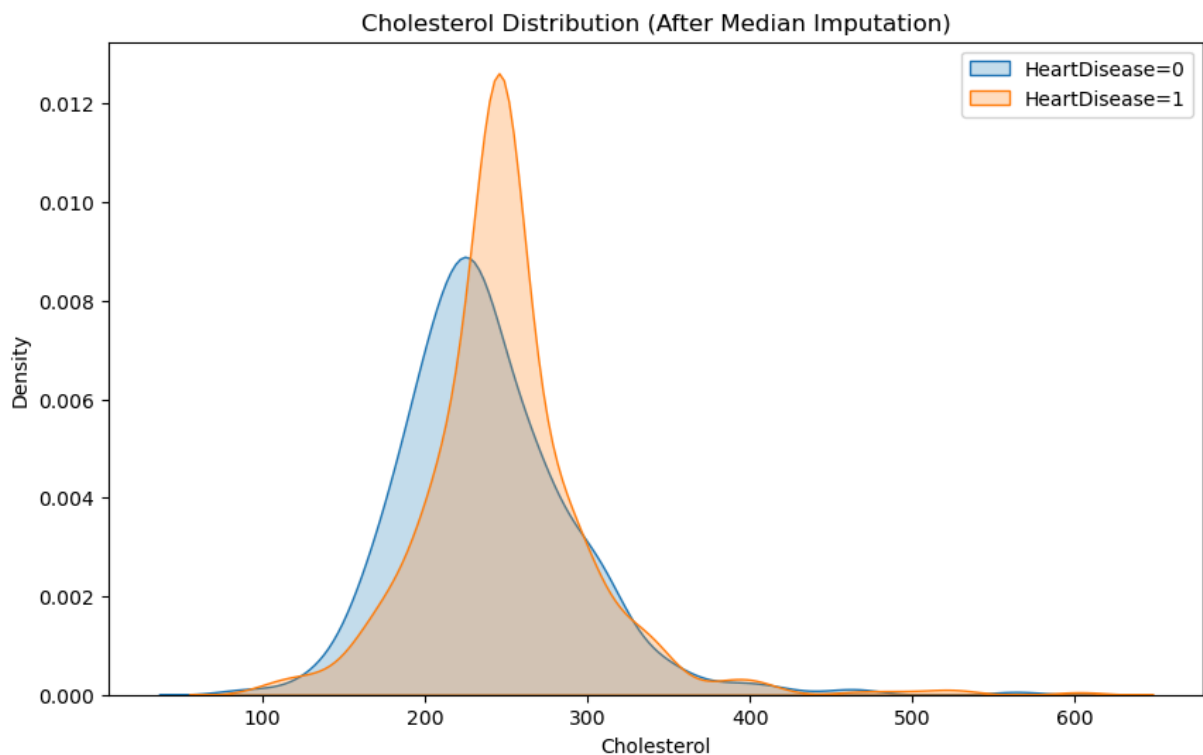
```

	count	mean	std	min	25%	50%	75%	m
ax								
HeartDisease								
0	410.0	238.414634	54.045994	85.0	204.0	231.5	266.75	56
4.0								
1	507.0	249.554241	52.370323	100.0	225.0	246.0	267.00	60
3.0								

```

In [72]: # KDE plot – distribution comparison after median imputation
plt.figure(figsize=(10,6))
sns.kdeplot(df_clean[df_clean["HeartDisease"]==0]["Cholesterol"], label="HeartDisease=0")
sns.kdeplot(df_clean[df_clean["HeartDisease"]==1]["Cholesterol"], label="HeartDisease=1")
plt.title("Cholesterol Distribution (After Median Imputation)")
plt.xlabel("Cholesterol")
plt.ylabel("Density")
plt.legend()
plt.show()

```



In [74]: `## Provides summary statistics for numeric columns for clean data, rounded to 2 decimal places`  
`df_clean.describe().round(2).T`

Out[74]:

	count	mean	std	min	25%	50%	75%	max
<b>Age</b>	917.0	53.51	9.44	28.0	47.0	54.0	60.0	77.0
<b>RestingBP</b>	917.0	132.54	18.00	80.0	120.0	130.0	140.0	200.0
<b>Cholesterol</b>	917.0	244.57	53.39	85.0	214.0	246.0	267.0	603.0
<b>FastingBS</b>	917.0	0.23	0.42	0.0	0.0	0.0	0.0	1.0
<b>MaxHR</b>	917.0	136.79	25.47	60.0	120.0	138.0	156.0	202.0
<b>Oldpeak</b>	917.0	0.89	1.07	-2.6	0.0	0.6	1.5	6.2
<b>HeartDisease</b>	917.0	0.55	0.50	0.0	0.0	1.0	1.0	1.0


In [76]: `# Count negative Oldpeak values`

```
neg_oldpeak_count = (df["Oldpeak"] < 0).sum()
print(f"Number of negative Oldpeak values: {neg_oldpeak_count}")
```

Number of negative Oldpeak values: 13

In [78]: `# Stripplot showing distribution of Oldpeak values by HeartDisease, with reference line at 0`

```
plt.figure(figsize=(8,5))
sns.stripplot(x="HeartDisease", y="Oldpeak", data=df, jitter=True, palette="magma")
plt.axhline(0, color="red", linestyle="--")
plt.title("Oldpeak Values by HeartDisease (individual points)")
plt.show()
```

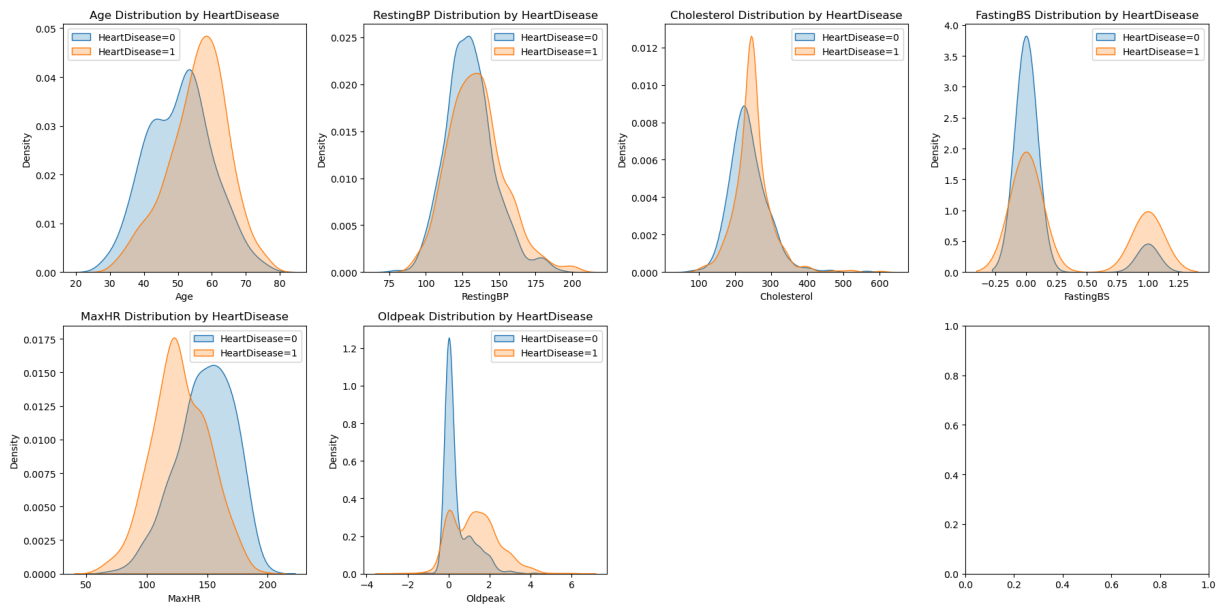


```
In [82]: # Select only numeric columns (excluding the target if needed)
num_cols = df_clean.select_dtypes(include="number").columns.drop("HeartDisease")

# Distributions of numerical features by HeartDisease status
fig, axes = plt.subplots(2, 4, figsize=(18, 9))
axes = axes.flatten()
for i, col in enumerate(num_cols):
    ax = axes[i]
    sns.kdeplot(df_clean[df_clean["HeartDisease"]==0][col].dropna(), label="0")
    sns.kdeplot(df_clean[df_clean["HeartDisease"]==1][col].dropna(), label="1")
    ax.set_title(f"{col} Distribution by HeartDisease"); ax.legend()

# Remove empty subplot (if number of cols < 8)
fig.delaxes(axes[len(num_cols)])
plt.tight_layout()
plt.show()
```





```
In [83]: # Select only numerical columns and to check correlation
num_cols = df_clean.select_dtypes(include=[np.number]).columns

plt.figure(figsize=(8,6))
sns.heatmap(df_clean[num_cols].corr(), annot=True, cmap="viridis", fmt=".2f")
plt.title("Correlation Heatmap (numeric features)")
plt.show()
```




```
In [84]: # Distribution of categorical variables by the target variable
cat_cols = ["Sex", "ChestPainType", "FastingBS", "RestingECG", "ExerciseAngi"]
fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(16, 14))
axes = axes.flatten()

for ax, col in zip(axes, cat_cols):
    g = sns.countplot(data=df_clean, x=col, hue="HeartDisease", palette="Set1")
    ax.set_title(f"{col} by HeartDisease")
    ax.set_xlabel(col); ax.set_ylabel("Count")
    # label name
    for c in g.containers:
        g.bar_label(c, padding=2, fmt="%.0f")

# Remove extra axes
for ax in axes[len(cat_cols):]:
    fig.delaxes(ax)

plt.tight_layout()
plt.show()
```

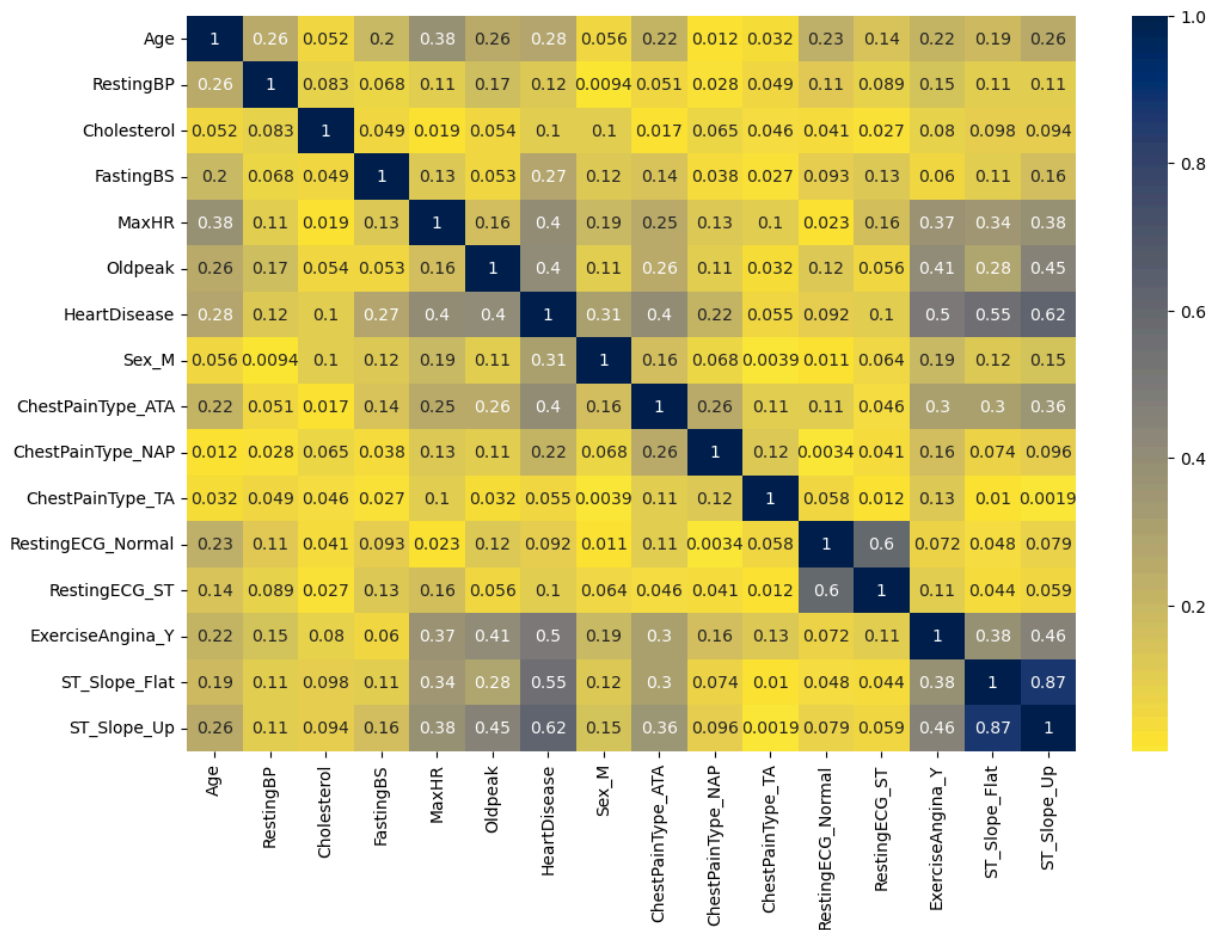


```
In [88]: # One-hot encoding was applied to transform categorical variables into dummy
DUMMY = pd.get_dummies(df_clean, drop_first=True)
DUMMY.head()
```

Out[88]:

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	HeartDisease	Sex_M
0	40	140	289.0	0	172	0.0	0	True
1	49	160	180.0	0	156	1.0	1	False
2	37	130	283.0	0	98	0.0	0	True
3	48	138	214.0	0	108	1.5	1	False
4	54	150	195.0	0	122	0.0	0	True

```
In [90]: ## Compute absolute pairwise correlations (after one-hot encoding) and visual
correlations = abs(DUMMY.corr())
plt.figure(figsize=(12,8))
sns.heatmap(correlations, annot=True, cmap="cividis_r")
plt.show()
```



In [ ]:

In [ ]: